PredictChain

Matthew Pisano, Connor Patterson, William Hawkins

Spring 2023

1 Project Description

1.1 Problem Solved

PredictChain helps to solve one of the main issues that involve AI models today: accessibility. Our project fulfils this need in two ways. Oftentimes, individuals or groups poses data that they would wish to be used in predictive analysis. However, these people may not have access to the compute capacity needed to train predictive models on this data. Additionally, yet other people have neither access to predictive training data, nor do they have access to computational resources. PredictChain solves both of these problems simultaneously.

When users upload their datasets, they allow a model to be trained on those datasets. Higher quality datasets will produce higher quality models. When users submit parameters for training, they allow the model that their parameters produce to be used publicly. Both of these users are rewarded for their work when a model is queried. The amount of this reward is based off of the correctness of the prediction. This encourages users to participate in contributing the resources needed for good predictions, while leaving a public record for other users to view.

1.2 Significance

As the trend of ever-growing machine learning models continues, this problem becomes increasingly relevant. As the scale and power of these models grow, the resources required to train them grow as well. This makes the training of useful machine learning models unattainable for most people. In order to get useful results from these models, users often have to pay large, centralized organizations, without any reward if they provide a good dataset or model parameters. PredictChain changes this paradigm by incentiveizing the thoughtful creation of useful models and datasets.

Additionally, after the recent mania and subsequent crash around blockchain adjacent technologies, it has become important to remind people that blockchains can be used for genuine utility in addition to investment. By primarily using Algos as a method of payment, instead of investment, it helps to, once again, show that cryptocurrencies can effectively be used as a pure form of payment for useful services. Of course, this is in addition to the many other services that use crypto in a similar manner, but adding one more project only helps the Algorand's notion of usefulness.

1.3 Use Cases and User Stories

User Story #1

Scenario: As a data analyst, I want to have access to trends of various stock markets so that I can create predictive models that will inform my investment strategies.

User Story #2

Scenario: As a stockbroker, I want to compare my dataset to another dataset while using the same model so that I can get an idea as to which is better.

Use Cases

For the sake of space, we will only go over two simple use cases; one that allows the user to add a dataset and another where a user queries a prediction from an existing model.

Table 1: Add Dataset

Identifier:	UC1
Description:	The user logs in to PredictChain and adds a dataset to their account
Actor(s):	Site User
Precondition(s):	The user has a publicly available link for the dataset
Event Flow:	1. The user logs in properly into their account (or creates a new account) 2. The user checks to see how much it would cost to upload a dataset, given the size of the dataset 3. The user inputs a link to the dataset file in the upload dataset area 4. The user inputs a name corresponding to that dataset 5. The user enters in the size of the dataset file in bytes 6. The user presses "Submit"
Postcondition(s):	1. A dataset has been created and is stored in the oracle 2. The transaction has been stored on the blockchain 3. The website displays the information onto the dashboard corresponding to the outcome

Table 2: Model Query

Identifier:	UC2
Description:	The user logs in to PredictChain and queries a prediction from an existing model
Actor(s):	Site User
Precondition(s):	The user is ready with information concerning the model training parameters
Event Flow:	1. The user logs in properly into their account (or creates a new account) 2. The user selects an existing model from the dropdown under the query model area 3. The user inputs a list of numbers corresponding to the input data of the model 4. The user submits their request using the submission button in that section
Postcondition(s):	1. The website displays the result of the query onto the dashboard 2. The result of the query is stored on the blockchain

2 Implementation Details

The structure of PredictChain is primarily broken up into two parts: the client and the oracle. Both of these parts interact with each other through the blockchain. The following diagram illustrates this relation:
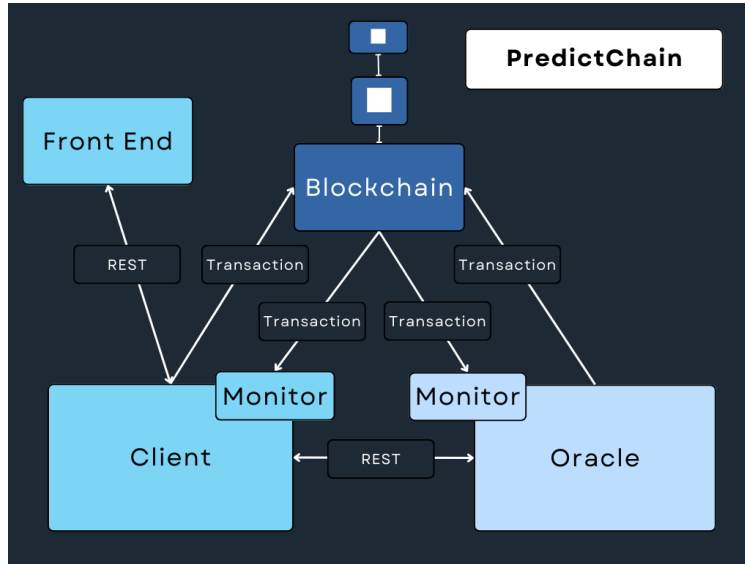

Figure 1: The architecture of PredictChain

The above diagram illustrates the core components of the system, along with their methods of communication.

2.1 The Front End UI

Although not directly related to the AI or the blockchain parts of the product, it is still important to discuss the UI aspect of our project. The UI sets the user's impression of PredictChain by having multiple pages that create a more pleasant experience for the user. For example, the home page talks about our mission, how we differ from our competitors and example model sets we provide. The addition of the *FAQ* and *Meet the Team* pages allow users to understand who created PredictChain, as well as get answers about any privacy/security concerns they might have. Lastly, users can create an account or login to a pre-existing account. Then users are able to use PredictChain to its fullest functionality such as adding datasets, checking prices, or querying models with provided datasets. The UI talks to the client through a series of REST requests in order to convey this data, later receiving the result of the user's actions later through the same means.

2.2 The Client

The client serves as a middleman between the front end user interface and the blockchain. It is run as a server, serving UI content to the user, taking in requests from the UI, and parsing those requests into a form suitable for both the blockchain and for the oracle. Additionally, the client constantly polls for updates coming from the oracle, through the blockchain, and then through its monitor. These updates are queued and sent to the front end upon request. This allows the user to both interact with the blockchain and to see the important updates that come from it.

2.3 The Oracle

Figure 2: The architecture of the Oracle

The oracle accomplishes the majority of the other tasks that this project requires. It constantly polls for updates coming from the client, through the blockchain by using its own monitor. Upon receiving these updates, it begins the execution of one of its three main operations. These are:

- Downloading a user-specified dataset and saving it

- Training one of the raw models based on user inputted parameters

- Querying one of the trained models on user inputted data and comparing it to the real-world result

After each of these operations, the oracle sends out several blockchain transactions. These can be either rewards to contributors of a model or confirmations/results of the operation that has been performed.

When working with user-submitted datasets, the oracle uses a handler to manage the operations performed on that dataset. The handler can save datasets to a specified environment, load datasets from a specified environment, parse that dataset as a pandas dataframe, and split the dataset by the values of one of its attributes. The environments that the handler recognizes are *local* and *IPFS*. When using either of these environments, the handler abstracts away the complexities of working with either of them into a unified interface.

When working with user-trained models, the oracle uses a similar, common interface. This interface can create the model architecture, train the model on a selected dataset, query the trained model, evaluate its performance, save the model, and load it back from a specified environment. When creating and training a model, the interface chooses among a group of archetype or template models. These models can be a:

- Multi-layered perceptron neural network

- Recurrent neural network

- Long short-term memory neural network

- Gated recurrent unit neural network

Each of these models has a *model_complexity* attribute. This is a simple float value, designed to give users a general idea of how performant a model can be once trained and serves as a method of calculating the cost of using or training that given model. The attribute itself is calculated using the size of the network and a linear multiplier to account for more complex model architectures. For models like GRUs or LSTMs, the complexity is higher as they are more complex and often better performing models. For models like MLPs, the complexity is lower. This gives the desired effect of faster, simpler models being cheaper than the slower, more complex models without any heavy calculations. The interface abstracts most of the complexities of training, querying, and evaluating these models. The only difference between them is the inclusion of several optional parameters.

2.4   The Blockchain

In PredictChain, the blockchain serves as both a records keeper and a messenger between the client and the oracle. This is accomplished by using transactions as a form of direct communication. With every transaction sent, there is a note. This note is a json-encoded string (encoded in base64) that communicates information about the operation that the transaction is requesting and arguments for that operation. These operations are represented by a series of op codes. These codes are abbreviations of the operation name enclosed in angle brackets, for example $\langle QUERY\_MODEL \rangle$. The arguments to these operations are represented as a named dictionary, with each key being the name of the argument and each value being the argument itself. This named strategy allows the program to be very flexible without worrying about the exact ordering of the arguments. Blockchain is quite useful in its role due to its immutability and its transparency. Using a blockchain means that all requests are permanently stored and public, so other users can see what type of models are useful for specific datasets and what results those models have produced.

2.5   Software and Libraries

Python Libraries

This project is primarily built in Python, using the Algorand SDK. The SDK makes interacting with the blockchain very straightforward. Through the tools provided by this library, we can easily read and write transactions from the Algorand blockchain.

Through Python, we also used data science and machine learning libraries such as Pandas and Torch. These libraries encapsulate many of the complexities of data preparation and model training for us. By using these libraries, we were able to concentrate on the higher-level functions of the project instead of worrying about the lower-level implementation.

Flask as also an important part of the project. We used Flask to allow both the client and oracle nodes to function as servers. The client would take in requests from the user and send out requests to the oracle. The oracle would then take in those requests and issue responses. We chose to use Flask in client-oracle communication to cut down on the amount of trivial transactions that would otherwise be made. For example, it would not benefit the accessibility or transparency of the project greatly if the exchanged transactions were dominated by simple *'what is the price of . . .'* requests.

Node Libraries

Additionally, the front end utilizes the React framework and firebase. React was useful to us as it streamlined the process of making dynamic, modular code for the web interface. Firebase was invaluable for handling administrative tasks such as keeping track of registered users and their associated metadata.

Redis

For the recommended configuration of the project, we use Redis as well. Redis helps to provide a reliable store for our metadata about models and datasets in a simple, persistent manner.

2.6 Resource Links

The following table will provide links to the various resources relevant to our project and its evaluation:

Table 3: Project Resources

Resource	Link
GitHub	github.com/AI-and-Blockchain/S23_PredictChain
Python Documentation	github.com/AI-and-Blockchain/S23_PredictChain/blob/main/docs/sphinx/index.html (Best viewed when opened locally with a web browser)
Overall Demo Video	drive.google.com/file/d/1vv3BEMC5ru3oa1HLSEXGNvjdFcZbCas3/view?resourcekey
Technical Demo Video	www.youtube.com/watch?v=icWc1qvhsgY

3 Evaluation

While the most definitive evaluation would be to deploy our project and get feedback from actual users, we are limited in our time and in our scope. In place of this, we have devised several tests that are designed to evaluate each component of the project on its own and how the entire project functions as a whole.

3.1 Transactions

The usage of transactions is central to the communications protocol of PredictChain, so making sure the protocol functions correctly is critical to the evaluation of the project. Thanks to the SDK, we had no issue with encoding the notes and actually sending the transactions. Where problems can potentially arise is within the client and oracle monitors.

The monitors are classes inside the client and oracle that listen for any transactions that have their node address as the recipient. This listening is done using the Algorand indexer class and a constant polling for new transactions. We noticed that this monitor would sometimes skip or duplicate incoming transactions. As we developed fixes for these issues, we constantly evaluated the performance of the monitor. This was done through programmatically sending one or more transactions to the client or oracle address and checking to see how the monitor handled them. By comparing the unique transaction ids of the sent transactions to those processed, we were able to evaluate the monitor, identify issues, and create fixes. For example, we now constantly update the minimum timestamp the indexer can look for transactions after and we also keep a registry of the ids of all past transactions. This helps to eliminate the duplicate transactions that were getting through.

3.2 Models

Another critical component of PredictChain is its usage of a variety of models. At present, all of these models are different types of neural networks. However, these networks are not created equally, with each having different qualities, architectures, and outputs. In order to evaluate the performance of these models, we ran several experiments on the models where their hyper-parameters and architectures were kept constant as they were tested. For this evaluation, we measured their performance on our sample dataset, The University of California, Irvine's *Dow Jones Index* data set [1]. This dataset has 16 different parameters, detailing the attributes of a range of stocks from the first half of 2011. For our usage, we do not eliminate any of these during training, although this may be an opportunity for future improvement. As for the models, they were all initialized with the following parameters:

Table 4: Model Evaluation Parameters

Parameter Name	Value	Description
epochs	70	The number of epochs that the model would train for
target_attrib	*close*	The attribute that the model was trying to predict, this is the daily closing price of the stock
hidden_dim	5	The number of neurons in each hidden layer
num_hidden_layers	1	The number of hidden layers
time_lag	0	The number of time steps that pass between the input window and the prediction
training_lookback	10	The number of time steps that recurrent models receive as input
sub_split_value	0	The integer id of the stock to predict, in this case its is *$AA* (Alcoa Corp)

We performed this test on all of our basic model structures, specifically our GRU, LSTM, RNN, and MLP models. In the following figures, we show the input data and predictions for each model. We also show the final loss and the final accuracy. The loss is calculated using the mean absolute error function $mae(x, y) = (\Sigma_{i=1}^{n} |y_i - x_i|)/n$ and our accuracy is calculated using $acc(x, y) = \sigma(-mae(x, y) + e^2)$ where σ is the sigmoid function. Our accuracy function is modified in this way so that very large losses are registered as somewhat accurate and lower losses are registered as very accurate. This helps to account for the very large losses generated by some of the less performant models.

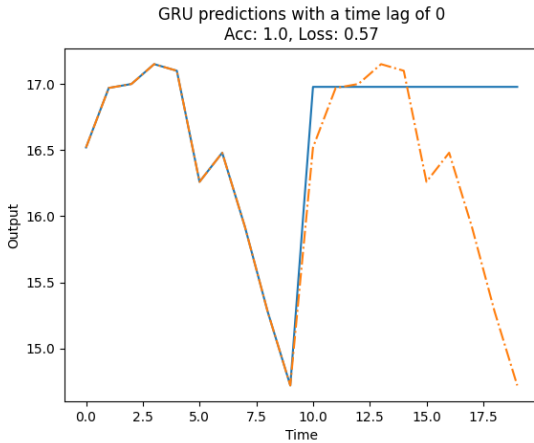The results of our evaluations are as follows:



Figure 3: The results from our GRU model

Figure 4: The results from our LSTM model

Figure 5: The results from our RNN model

Figure 6: The results from our MLP model

Each of these results reflects the individual strengths and weaknesses of each model. As is reflected in our *model_complexity* calculations, the GRU model is the most performant, with a loss of only 0.57. This is closely followed by the LSTM model, which still offers a good prediction, but is slightly less consistent overall. In our evaluation of our RNNs, we noticed a good example of the exploding gradients problem. Without the limits provided by the GRU and LSTM models, the predictions of the RNN become increasingly erratic. Finally, our MLP model does not suffer from an exploding gradient, but it does lack the recurrence of the other three models, only being able to output one day at a time.

3.3 End-To-End Testing

The final component of our testing was the end-to-end series of tests we performed. The goal of these tests was to give us a holistic picture of how well our project functioned and how well that functioning met our stated goals. For all of these tests, we started with front end user interaction and ended with the response to those actions.

Price Queries   The first, and most straightforward, set of operations that we tested were the three price query operations. We performed this test by submitting query requests on the front end, then tracing the functions called by that operation from the client to the oracle, then to the oracle getting the price, then finally back to the client. We performed this test for all three of our queries, testing different sets of inputs, both valid and invalid, seeing if they were handled properly.

Major Transactions   The next set of tests we performed focused on our major transactions: *Upload Dataset*, *Train Model*, and *Query Model*. These tests were similar to those for the price queries. As we performed these tests, we noted that, while the operations did complete successfully, the user feedback and results were ambiguous. Additionally, we realized that it may be confusing for the user to remember the exact names of the models and datasets in order to use them. To address these issues, we made several additions. To address the issue of users having to input the exact names, we added a dropdown feature instead of a test input. We now store a list of all current datasets and models in the system. This is updated upon the reloading of the page or when a user submits a transaction. This ensures the list is always updated for better ease of use. To address the feedback issue, we added a feedback section that periodically pinged the client to see if any new response transactions came in from the oracle. If these transactions did come in, we would display both the operation, the model or dataset name, and any extra data (like the result of a query) to the user. This way, the user would not have to manually check the transaction note on the Algorand block explorer.

By performing these tests, we were able to both evaluate the quality of the project as a whole and make important additions to improve the user experience. Our transaction tests helped us to identify the issues that had previously existed and to verify that the current communication protocol was working properly. Our model tests helped us to confirm our previous assumptions about the nature of our various models and which scenarios they are most useful in. Finally, our end-to-end testing confirmed the overall functioning of the project and inspired us to make some valuable improvements to the user experience.

4 Conclusion

4.1 Summary

Over the course of our work, we have made PredictChain into a working blockchain-based marketplace for predictive AI models. Through PredictChain, users are now able to upload datasets for training predictive models, request that basic models be trained on any previously uploaded datasets, or submit queries to those trained models. These various models will be operated by a central node with computing resources available. A variety of models are available, ranging from cheap, fast, and simple to more expensive, slower, and more powerful. This will allow for a large variety of predictive abilities for both simple and complex patterns. All the past predictions form these models will be stored on the blockchain for public viewing.

4.2 Limitations

A notable limitation that we experienced during this project were constraints on both our time and the amount of hours we could each put into the project. If this project was worked on for several months longer, the final product would be significantly more fleshed out with more features. If this were more akin to a multi-semester project, we would have had the opportunity to add more starting datasets and add a more diverse set of models.

Another limitation that we had experienced was that we each had other commitments to other classes or activities. This limited the amount of hours per week we could each work on improving the project. To compensate for this, we worked together to dynamically adjust our efforts by picking up the work of anyone that was unable to put as much time into the project. This strategy helped us to stay on track while maintaining our prior commitments.

Despite, or perhaps aided by, these limitations, were able to budget our time and distribute the work in such a way that out final project is in a functioning state that fulfils our initial objective: to create a prototype for a accessible and transparent marketplace for machine learning training and predictions.

4.3 Future Work

As mentioned above, we have several opportunities for improvement that could be accomplished with future work into this project. One major improvement that we could make to the project is the overhaul of the architecture of the project. As mentioned previously, we had briefly considered using many model training nodes instead of a centralized node. Adding this would make the project more decentralized and open up another opportunity for community contribution. This would likely come in the form of users hosting training nodes and being rewarded for the quality of models that it produces. Another potential improvement that could be introduced with future work is the addition of a greater variety of models or more example datasets. Currently we only use neural networks to use for predictions. In future work, we would like to add a more diverse set of models, such as decision trees or more statistical models based off of Bayesian inference. Finally, we could make improvements to how our models are trained, for instance, allowing users to prune off attributes from a dataset that they did not deem useful to the trained model. Adding these improvements may make PredictChain a more fleshed-out version of itself without changing the core principles of our project.

References

[1] U. of California Irvine. Dow jones index data set. https://archive.ics.uci.edu/ml/datasets/Dow+Jones+Index.